

Big Data for Public Policy

3. Machine Learning Essentials

Elliott Ash & Malka Guillot

Where we are


- ▶ Past weeks:
 - ▶ w1: Overview and motivation
 - ▶ w2: Finding datasets using webcrawling and API
- ▶ This week (w3): Intro to supervised Machine Learning (ML) - regressions
- ▶ Next:
 - ▶ w4: Text analysis fundamentals
 - ▶ w5: Supervised Machine Learning (ML) - classification
 - ▶ w6: Unsupervised Machine Learning (ML)

Today: supervised ML - regressions

- ▶ First hour:
 - ▶ What is machine learning?
 - ▶ Basic steps and concepts
- ▶ Second hour:

Application Predicting the prices of houses using given features

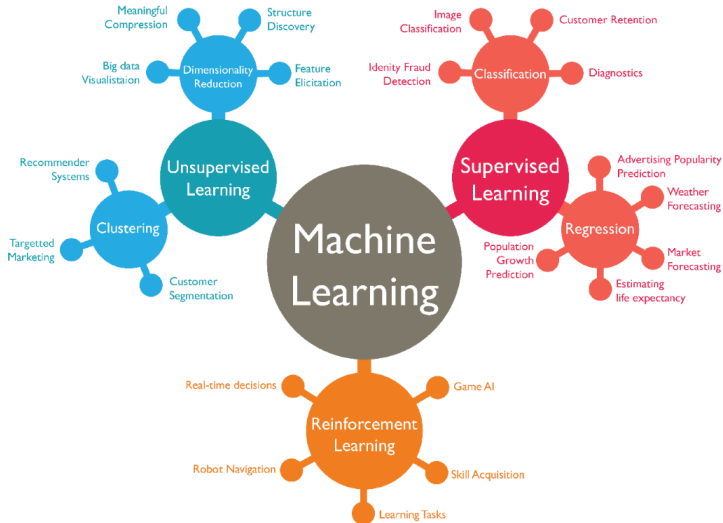
What is machine learning?



- ▶ In classical computer programming, humans input the rules and the data, and the computer provides answers.
- ▶ In machine learning, humans input the data and the answers, and the computer learns the rules.

⇒ **Machine learning is the science (and art) of programming computers so they can *learn from data*.**

The Machine Learning Landscape



Machine Learning

Usually, ML is divided in 2 categories:

- ▶ the predictive or supervised learning approach
- ▶ the descriptive or unsupervised learning approach

Econometrics vs. Machine Learning

- ▶ In **econometrics**, we estimate a low-dimensional **causal parameter** θ .
 - ▶ goal of estimation: unbiasedness
 - ▶ predicts how outcome y would change if treatment variable x were **exogenously shifted**.
 - ▶ useful for policy evaluation.

Econometrics vs. Machine Learning

- ▶ In **econometrics**, we estimate a low-dimensional **causal parameter** θ .
 - ▶ goal of estimation: unbiasedness
 - ▶ predicts how outcome y would change if treatment variable x were **exogenously shifted**.
 - ▶ useful for policy evaluation.
- ▶ In **machine learning**, we approximate a function $h(X; \theta)$ to predict Y given covariates X .
 - ▶ goal of prediction: loss minimization
 - ▶ if we collected more data on X , we could predict the associated \hat{Y} .

Econometrics vs. Machine Learning

- ▶ In **econometrics**, we estimate a low-dimensional **causal parameter** θ .
 - ▶ goal of estimation: unbiasedness
 - ▶ predicts how outcome y would change if treatment variable x were **exogenously shifted**.
 - ▶ useful for policy evaluation.
- ▶ In **machine learning**, we approximate a function $h(X; \theta)$ to predict Y given covariates X .
 - ▶ goal of prediction: loss minimization
 - ▶ if we collected more data on X , we could predict the associated \hat{Y} .
 - ▶ but $h(\cdot)$ does not provide a *counterfactual prediction* – that is, how the outcome would change if X 's were exogenously shifted.

Example of Applications

- ▶ Detect fraud: taxes, social benefits
- ▶ Forecasts next year's revenue
- ▶ Diagnosis of diseases

Outline

Machine Learning Overview

Overfitting and Regularization

Pipelines and Cross-Validation

Basic Setup

- ▶ Suppose we have m observations within a dataset of the form (y_i, x_i) for $i = 1, \dots, m$:
 - ▶ y_i : dependent / response / label variable
 - ▶ x_i : P -dimensional vector of independent variables, covariates or features. Potentially, $P \gg N$ it's possible to have more variables than observations
- ▶ Supervised Learning: Learn a mapping from x_i to y_i
 - ▶ Classification problem: y_i is categorical
 - ▶ Regression problem: y_i is continuous
- ▶ Unsupervised Learning: Learn some structure within the x_i observations → Description problem

What do ML Algorithms do? Minimize a cost function

- ▶ A typical cost function for regression problems is Mean Squared Error (MSE):

$$\text{MSE}(X, h) = \frac{1}{m} \sum_{i=1}^m (h(x_i) - y_i)^2$$

- ▶ m , the number of rows/observations
- ▶ X , the feature set, with row x_i
- ▶ Y , the outcome, with item y_i
- ▶ $h(x_i)$ the model prediction (hypothesis)

Linear Regression is Machine Learning

- ▶ OLS assumes the functional form

$$y_i = x_i' \theta + \epsilon_i$$

and minimizes the MSE

$$\min_{\hat{\theta}} \frac{1}{m} \sum_{i=1}^m (x_i' \hat{\theta} - y_i)^2$$

Linear Regression is Machine Learning

- ▶ OLS assumes the functional form

$$y_i = x_i' \theta + \epsilon_i$$

and minimizes the MSE

$$\min_{\hat{\theta}} \frac{1}{m} \sum_{i=1}^m (x_i' \hat{\theta} - y_i)^2$$

- ▶ This has a closed form solution

$$\hat{\theta} = (\mathbf{X}' \mathbf{X})^{-1} \mathbf{X}' \mathbf{Y}$$

- ▶ But

- ▶ this solution does not work well with large m and n
- ▶ most machine learning models do **not** have a closed form solution.

How do ML Algorithms Work? Gradient Descent




Figure 4-3. Gradient Descent

- ▶ Gradient descent measures the local gradient of the error function, and then steps in that direction.
 - ▶ Once the gradient equals zero, you have reached a minimum.

Evaluating Algorithms: Cross Validation

Training and Testing

- ▶ Machine learning models can achieve arbitrarily high accuracy in-sample.
 - ▶ performance should be evaluated out-of-sample.

Training and Testing

- ▶ Machine learning models can achieve arbitrarily high accuracy in-sample.
 - ▶ performance should be evaluated out-of-sample.
- ▶ Standard approach: Train-Set Sample Split
 - ▶ randomly sample 80% of the data as a training dataset
 - ▶ data transformations, feature selection, and hyperparameter tuning only see this data.

Training and Testing

- ▶ Machine learning models can achieve arbitrarily high accuracy in-sample.
 - ▶ performance should be evaluated out-of-sample.
- ▶ Standard approach: Train-Set Sample Split
 - ▶ randomly sample 80% of the data as a training dataset
 - ▶ data transformations, feature selection, and hyperparameter tuning only see this data.
 - ▶ form predictions in 20% test dataset to evaluate performance.

Outline

Machine Learning Overview

Overfitting and Regularization

Pipelines and Cross-Validation

The Problem of Overfitting




Figure 4-14. High-degree Polynomial Regression

Bias-Variance Tradeoff

- ▶ *Bias*
 - ▶ Error due to wrong assumptions, such as assuming data is linear when it is quadratic.
 - ▶ Will underfit the training data

Bias-Variance Tradeoff

- ▶ *Bias*

- ▶ Error due to wrong assumptions, such as assuming data is linear when it is quadratic.
- ▶ Will underfit the training data

- ▶ *Variance*

- ▶ Error due to excess sensitivity to small variations in the training data.
- ▶ A model with high variance is likely to overfit the training data.

Bias-Variance Tradeoff

- ▶ *Bias*

- ▶ Error due to wrong assumptions, such as assuming data is linear when it is quadratic.
- ▶ Will underfit the training data

- ▶ *Variance*

- ▶ Error due to excess sensitivity to small variations in the training data.
- ▶ A model with high variance is likely to overfit the training data.

- ▶ *Irreducible error*

- ▶ Error due to noise in the data.

Training Notes

- ▶ In general, increasing a model's complexity will increase variance and reduce bias.

Training Notes

- ▶ In general, increasing a model's complexity will increase variance and reduce bias.
- ▶ If the model is underfitting:
 - ▶ adding more training data will not help.
 - ▶ use a more complex model

Training Notes

- ▶ In general, increasing a model's complexity will increase variance and reduce bias.
- ▶ If the model is underfitting:
 - ▶ adding more training data will not help.
 - ▶ use a more complex model
- ▶ If the model is overfitting:
 - ▶ adding more training data may help
 - ▶ or use regularization
 - ▶ cross-validation

Univariate OLS for Ranking Predictive Features

- ▶ Consider the univariate regression

$$Y_i = \theta_w x_i^w + \epsilon_i$$

for each text feature w (e.g., relative word or n-gram frequency).

- ▶ Can be estimated with OLS.

Univariate OLS for Ranking Predictive Features

- ▶ Consider the univariate regression

$$Y_i = \theta_w x_i^w + \epsilon_i$$

for each text feature w (e.g., relative word or n-gram frequency).

- ▶ Can be estimated with OLS.
 - ▶ Can add fixed effects:

$$Y_i = \alpha_i + \theta_w x_i^w + \epsilon_i$$

where α_i is a vector of dummy variables (fixed effects) for categories, e.g. location, year.

Univariate OLS for Ranking Predictive Features

- ▶ Consider the univariate regression

$$Y_i = \theta_w x_i^w + \epsilon_i$$

for each text feature w (e.g., relative word or n-gram frequency).

- ▶ Can be estimated with OLS.
 - ▶ Can add fixed effects:

$$Y_i = \alpha_i + \theta_w x_i^w + \epsilon_i$$

where α_i is a vector of dummy variables (fixed effects) for categories, e.g. location, year.

- ▶ Even better, can residualize Y and X on fixed effects before running any regressions.
 - ▶ That is, regress $Y_i = \alpha_i + \epsilon_i$ and $x_i^w = \alpha_i + \epsilon_i, \forall w$, take residuals $\tilde{Y}_i = Y_i - \hat{\alpha}_i$ and $\tilde{x}_i^w = x_i^w - \hat{\alpha}_i$

Univariate OLS for Ranking Predictive Features

- ▶ Consider the univariate regression

$$Y_i = \theta_w x_i^w + \epsilon_i$$

for each text feature w (e.g., relative word or n-gram frequency).

- ▶ Can be estimated with OLS.

- ▶ Can add fixed effects:

$$Y_i = \alpha_i + \theta_w x_i^w + \epsilon_i$$

where α_i is a vector of dummy variables (fixed effects) for categories, e.g. location, year.

- ▶ Even better, can residualize Y and X on fixed effects before running any regressions.

- ▶ That is, regress $Y_i = \alpha_i + \epsilon_i$ and $x_i^w = \alpha_i + \epsilon_i, \forall w$, take residuals $\tilde{Y}_i = Y_i - \hat{\alpha}_i$ and $\tilde{x}_i^w = x_i^w - \hat{\alpha}_i$

- ▶ then regress

$$\tilde{Y}_i = \theta_w \tilde{x}_i^w + \epsilon_i$$

- ▶ This can be used for descriptives, and also for feature selection

- ▶ For quick feature selection, can use sklearn's `f_regression`

Ridge, Lasso, and Elastic Net

- ▶ Ridge and lasso regression are tools for dealing with large feature sets where:
 - ▶ models have multicollinearity that causes bias
 - ▶ models tend to overfit
 - ▶ models are computationally costly to fit
- ▶ These algorithms work by constraining estimated parameter sizes.

Ridge Regression

- ▶ The Ridge cost function is

$$J(\theta) = \text{MSE}(\theta) + \underbrace{\alpha_2 \frac{1}{2} \sum_{i=1}^n \theta_i^2}_{\text{Regularization term}}$$

- ▶ i indexes over n features
- ▶ α_2 is a **hyperparameter** setting the strength of the L2 penalty
 - theta's for all the model features

Ridge Regression

- ▶ The Ridge cost function is

$$J(\theta) = \text{MSE}(\theta) + \underbrace{\alpha_2 \frac{1}{2} \sum_{i=1}^n \theta_i^2}_{\text{Regularization term}}$$

- ▶ i indexes over n features
 - ▶ α_2 is a **hyperparameter** setting the strength of the L2 penalty
- ▶ Ridge penalizes large coefficients, which reduces over-fitting to the training set.
 - ▶ The estimated coefficients, when taken to other data, will generalize better.

Ridge Regression

- ▶ The Ridge cost function is

$$J(\theta) = \text{MSE}(\theta) + \underbrace{\alpha_2 \frac{1}{2} \sum_{i=1}^n \theta_i^2}_{\text{Regularization term}}$$

- ▶ i indexes over n features
- ▶ α_2 is a **hyperparameter** setting the strength of the L2 penalty
- ▶ Ridge penalizes large coefficients, which reduces over-fitting to the training set.
 - ▶ The estimated coefficients, when taken to other data, will generalize better.
- ▶ It turns out that the Ridge estimator, like OLS, has a closed-form solution:

$$\hat{\theta}_{\text{Ridge}} = (X'X + \alpha_2 \mathbf{I}_n)^{-1} X' \mathbf{y}$$

where \mathbf{I}_n is the identity matrix.

- ▶ But it can also be solved by (stochastic) gradient descent.

Lasso Regression

- ▶ Least Absolute Shrinkage and Selection Operator Regression
- ▶ The Lasso cost function is

$$J(\theta) = \text{MSE}(\theta) + \alpha_1 \sum_{i=1}^n |\theta_i|$$

- ▶ i indexes over n features
- ▶ α_1 is a hyperparameter setting the strength of the L1 penalty

Lasso Regression

- ▶ *Least Absolute Shrinkage and Selection Operator Regression*
- ▶ The Lasso cost function is

$$J(\theta) = \text{MSE}(\theta) + \alpha_1 \sum_{i=1}^n |\theta_i|$$

- ▶ i indexes over n features
 - ▶ α_1 is a hyperparameter setting the strength of the L1 penalty
- ▶ Lasso automatically performs feature selection and outputs a sparse model.
- ▶ It does not have a closed-form solution but can be solved by gradient descent.

Elastic Net

- ▶ Elastic Net uses both L1 and L2 penalties:

$$J(\theta) = \text{MSE}(\theta) + \alpha_1 \sum_{i=1}^n |\theta_i| + \alpha_2 \frac{1}{2} \sum_{i=1}^n \theta_i^2$$

- ▶ in general, elastic net is preferred to lasso, which can behave erratically when the number of features is greater than the number of rows, or when some features are highly collinear.
 - ▶ but you have to tune two hyperparameters rather than one

Hyperparameters vs. Parameters

- ▶ **Parameters:** internal to the model whose values can be estimated from the data and we are often trying to estimate them as best as possible
 - ▶ **hyperparameters:** external to the model and cannot be directly learned from the regular training process
- ⇒ model-specific properties that are *fixed* before the model is trained

Hyperparameters tuning

- ▶ Use `GridSearchCV` or `RandomizedSearchCV` to automate search over parameter space.
 - ▶ For example: Elastic net hyperparameters should be selected to optimize out-of-sample fit.
 - ▶ “Grid search” scans over the hyperparameter space $(\alpha_1 \geq 0, \alpha_2 \geq 0)$, computes out-of-sample MSE for all pairs (α_1, α_2) , and selects the MSE-minimizing model.

Regularized Models Require Standardized Data

- ▶ Regularized models are designed to work with standardized predictors:

$$\tilde{x}_i = \frac{x_i - \bar{x}}{SD[\mathbf{x}]}$$

Regularized Models Require Standardized Data

- ▶ Regularized models are designed to work with standardized predictors:

$$\tilde{x}_i = \frac{x_i - \bar{x}}{SD[\mathbf{x}]}$$

- ▶ Taking out the mean will convert sparse data to dense data, you can avoid that by just dividing by the standard deviation:

$$\tilde{x}_i = \frac{x_i}{SD[\mathbf{x}]}$$

- ▶ in sklearn, set `with_mean=False`.

Outline

Machine Learning Overview

Overfitting and Regularization

Pipelines and Cross-Validation

Data Prep for Machine Learning

- ▶ See Geron Chapter 2 for pandas and sklearn syntax:
 - ▶ imputing missing values.
 - ▶ feature scaling (often helpful/necessary for ML models to work well)
 - ▶ encoding categorical variables.
 - ▶ see jupyter notebook
- ▶ Best practice: **reproducible data pipeline.**

Data Prep for Machine Learning

- ▶ See Geron Chapter 2 for pandas and sklearn syntax:
 - ▶ imputing missing values.
 - ▶ feature scaling (often helpful/necessary for ML models to work well)
 - ▶ encoding categorical variables.
 - ▶ see jupyter notebook
- ▶ Best practice: **reproducible data pipeline**.
- ▶ **Key point**: all data transformations, feature selection, and hyperparameter tuning must be done **in the training set**.

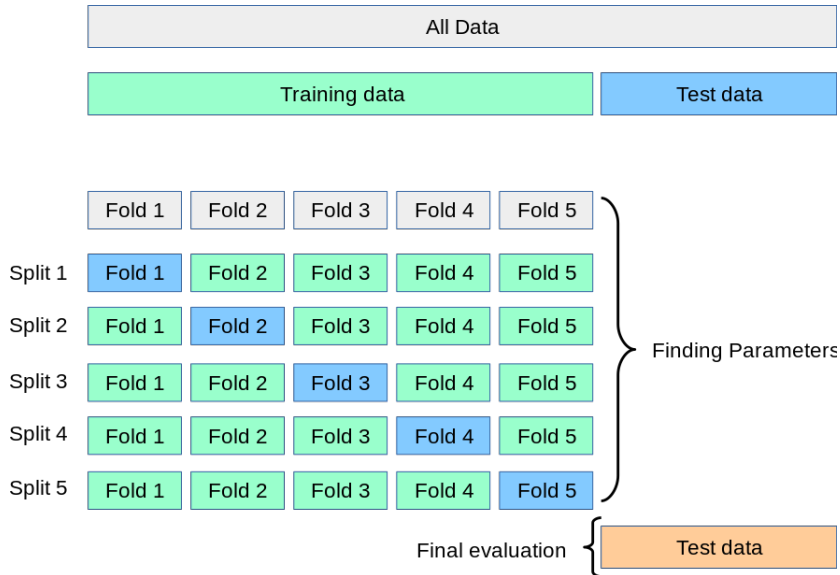
Cross-Validation

- ▶ Use `cross_val_score` method to get model performance across subsets of the training set:
 - ▶ split data into K folds.
 - ▶ for each fold $k \in \{1, 2, \dots, K\}$, train model in rest of data ($-k$) and evaluate MSE in k .
 - ▶ Report mean and s.d. of MSE across folds.

Cross-Validation

- ▶ Use `cross_val_score` method to get model performance across subsets of the training set:
 - ▶ split data into K folds.
 - ▶ for each fold $k \in \{1, 2, \dots, K\}$, train model in rest of data ($-k$) and evaluate MSE in k .
 - ▶ Report mean and s.d. of MSE across folds.

Cross-Validation




Evaluating Regression Models: R^2

- ▶ Mean squared error (MSE) can be used to compare regression models, but the units depend on the outcome variable and therefore are not interpretable.
 - ▶ Use R^2 in the test set for a more interpretable evaluation metric.
 - ▶ MSE provides same ranking as R^2 and is faster to compute, so it's preferred for model selection.

Evaluating Regression Models: Binscatter Plots

- Binscatters provide visual evidence of regression model performance: Plot Y against \hat{Y} in the test set:



Sample from Ash and Labzina (2019). $R^2 = .03$.

A Machine Learning Project, End-to-End

Aurelien Geron, *Hands-on machine learning with Scikit-Learn, Keras, & TensorFlow*, Chapter 2:

1. Look at the big picture.
2. Get the data.
3. Discover and visualize the data to gain insights.
4. Prepare the data for Machine Learning algorithms.

A Machine Learning Project, End-to-End

Aurelien Geron, *Hands-on machine learning with Scikit-Learn, Keras, & TensorFlow*, Chapter 2:

1. Look at the big picture.
2. Get the data.
3. Discover and visualize the data to gain insights.
4. Prepare the data for Machine Learning algorithms.
5. Select a model and train it.
6. Fine-tune your model:
 - 6.1 Testing and validating
 - 6.2 Hyperparameter tuning and model selection.
7. Present your solution.
8. Launch, monitor, and maintain your system.

ML with Python

- ▶ Python packages
 - ▶ pandas for handling tabular data
 - ▶ matplotlib and seaborn for plotting
 - ▶ Scikit-Learn (sklearn) for standard (non-deep-learning) ML algorithm
 - ▶ scikit-learn is very comprehensive and the online-documentation itself provides a good introduction into ML.
 - ▶ TensorFlow, PyTorch and Keras for deep-learning.

Scikit-Learn Design Principles

- ▶ **Consistency:**

- ▶ *Estimator*: An object that can estimate parameters. Estimation is performed by `fit()` method. Exogenous parameters (provided by the researcher) are called *hyperparameters*.

Scikit-Learn Design Principles

► Consistency:

- *Estimator*: An object that can estimate parameters. Estimation is performed by `fit()` method. Exogenous parameters (provided by the researcher) are called *hyperparameters*.
- *Transformer*: An object that transforms a data set. Transformation is performed by the `transform()` method. The convenience method `fit_transform()` both fits an estimator and returns the transformed input data set.

Scikit-Learn Design Principles

► Consistency:

- *Estimator*: An object that can estimate parameters. Estimation is performed by `fit()` method. Exogenous parameters (provided by the researcher) are called *hyperparameters*.
- *Transformer*: An object that transforms a data set. Transformation is performed by the `transform()` method. The convenience method `fit_transform()` both fits an estimator and returns the transformed input data set.
- *Predictor*: An object that forms a prediction from an input data set. The `predict()` method forms the predictions. It also has a `score()` method that measures the quality of the predictions given a test set.

Scikit-Learn Design Principles

► **Consistency:**

- *Estimator*: An object that can estimate parameters. Estimation is performed by `fit()` method. Exogenous parameters (provided by the researcher) are called *hyperparameters*.
- *Transformer*: An object that transforms a data set. Transformation is performed by the `transform()` method. The convenience method `fit_transform()` both fits an estimator and returns the transformed input data set.
- *Predictor*: An object that forms a prediction from an input data set. The `predict()` method forms the predictions. It also has a `score()` method that measures the quality of the predictions given a test set.

- **Inspection**: Hyperparameters and parameters are accessible. Learned parameters have an underscore suffix (e.g. `lin_reg.coef_`)

Scikit-Learn Design Principles

► **Consistency:**

- *Estimator*: An object that can estimate parameters. Estimation is performed by `fit()` method. Exogenous parameters (provided by the researcher) are called *hyperparameters*.
- *Transformer*: An object that transforms a data set. Transformation is performed by the `transform()` method. The convenience method `fit_transform()` both fits an estimator and returns the transformed input data set.
- *Predictor*: An object that forms a prediction from an input data set. The `predict()` method forms the predictions. It also has a `score()` method that measures the quality of the predictions given a test set.

► **Inspection:** Hyperparameters and parameters are accessible. Learned parameters have an underscore suffix (e.g. `lin_reg.coef_`)

► **Non-proliferation of classes:** Use native Python data types; existing building blocks are used as much as possible.

Scikit-Learn Design Principles

► **Consistency:**

- *Estimator*: An object that can estimate parameters. Estimation is performed by `fit()` method. Exogenous parameters (provided by the researcher) are called *hyperparameters*.
- *Transformer*: An object that transforms a data set. Transformation is performed by the `transform()` method. The convenience method `fit_transform()` both fits an estimator and returns the transformed input data set.
- *Predictor*: An object that forms a prediction from an input data set. The `predict()` method forms the predictions. It also has a `score()` method that measures the quality of the predictions given a test set.

- **Inspection**: Hyperparameters and parameters are accessible. Learned parameters have an underscore suffix (e.g. `lin_reg.coef_`)

- **Non-proliferation of classes**: Use native Python data types; existing building blocks are used as much as possible.

- **Sensible defaults**: Provides reasonable default values for hyperparameters – easy to get a good baseline up and running.